

Generating Slides from HTML pages

Arun JVS
IIIT-H

firstauthor@il.org

Rohit Girdhar
IIIT-H

secondauthor@i2.org

Sudheer Kumar
IIIT-H

secondauthor@i2.org

Abstract

some abstract TODO

slideshow is created which contains an summarized version of the paper. This Slideshow will contain the important points, tables, graphs, figures which helps to explain the paper. This may not serve as a final presentation but provides a good starting point for preparing it.

3. Related Work

There has been limited work that addresses the problem of automatic slides generation. Previous works include approaches like section-wise summarization [2], alignment methods for matching document regions with presentation regions [1], extracting topics and itemized elaborations from tagged documents. Many of these systems use multiple ideas from Information Retrieval, such as TF-IDF weights, query expansion, query-specific summarization, POS-tagging, etc.. Some of them are purely generative, while others learn models from an existing corpus of document-presentation pairs. These systems are tuned to work with different input formats such as PDF, XML, \LaTeX and PPT, each of which preserve a varying amount of semantic and structural information about the original text. In the following two subsections, we review two papers [2],[1] which particularly deal with technical papers.

3.1. SlidesGen

The first work we review is by Sravanthi et al. [2]. They propose an novel framework for automatic generation of presentation slides for technical papers. They take \LaTeX documents of research papers as input, and return the presentation slides. Their method depends on the assumption that by and large, conference papers have a similar structure: an abstract, followed by sections that can be broadly classified as introduction, related work, actual work (model), experiments, conclusion/results and bibliography. A slide in their system contains a title and some bulleted points that are important in that section. They evaluate the system by surveying the response of people who use their system.

Their system is divided into multiple stages. The first stage is pre-processing stage. In this step, the \LaTeX documents are converted into XML using a public domain con-

1. Motivation

Presenting information using slides is an effective way for audience retention. Presenters take the aid of slides for presenting research papers in conferences, for explaining them in classrooms and for other academic purposes. But creating slides from the research papers is a time consuming task for the presenter. So, it would be a great benefit for them if there is an easy way for generating the slides from the research papers. Also, sometimes people want to go through the research papers just to understand the overview of them. Slides provide an easier way for them to understand the overview of the topic. This gives rise to need for a tool which automatically generates slides from research papers. Automatic generation of slides from research paper is an easier task compared to generation of slides from some random document because research papers have an almost similar structure and we can address the problem of automatic generation of slides from research papers by exploiting the structure of the research papers. Once this problem is addressed, this work can be easily extended to automatically generate slides from other structured entities like book chapters etc.

2. Problem Definition

Given an HTML version of a research paper which is in accordance with the conference/journal proceedings, a

verter LaTeXML.

Next, they generate what they call "configuration file", which contains configuration parameters for each section (since each section has a different point of view and writing style). This involves categorization of the section and extraction of key phrases from the section which are then stored in the configuration file. For example, a section with large number of cite tags, or with title containing words such as "related work" or "literature survey" is categorized as related works section.

The next step is of extracting key phrases. Most research papers have associated key phrases that can help categorize the content in paper, and contain important concepts introduced in the paper. They are mostly related to model and experiments section, and can be used to summarize the same. So, the keywords given at the beginning are added to keyphrases for those sections in the configuration file. Also, few other phrases as the names of subsections etc are also added to the configuration file.

Next, they use QueSTS summarizer to summarize the model, experiment and conclusions section. QueSTS represents the text as a Integrated Graph where sentence is a node, and edge exists between 2 sentences if the sentences are similar. The edge weight is defined as the cosine similarity between the 2 sentences (above a given minimum threshold) Given the keyphrases computed previously, they are tokenized and a centrality based (to the query phrase) node weight is computed for the node corresponding to each token. Thus, they get query specific node weights and query independent edge weights.

From the above graph, they construct a Contextual Tree for each term q and node r . All the trees at node r are merged into a Summary Graph at r . The SGraphs generated from each node are ranked using a scoring model, and the one with best rank is returned as the summary.

The final step of this procedure is slide generation from the XML and configuration file. The introduction slides are generated by comparing introduction sentence with the abstract using cosine similarity and ones with high similarity are placed in the slides. They give a similar method for generating slides from related work section as well. From model and experiments section, the slides are generated using the above QueSTS mechanism. Similarly, conclusion slides are generated using comparison with keywords such as "proposed", "concluded" etc.

Another important aspect in slides is graphics. Graphics are added along with sentences that either refer to it or is present along with it.

The paper finally discusses the issues in alignment of sentences and generation of slides. They also evaluate their slides manually with multiple users, taking their feedback. Most users gave a satisfaction level of more than 8/10 to the slides generated.

3.2. Alignment Methods

This paper by Brandon et. al. [1] takes a different approach to slide generation, inspired by the human thought process while making a slide out of a paper. They focus their efforts on *alignment* methods - first breaking up the document and presentation into regions and then performing a matching on them. *Slide regions* include bullets, headings, and other text spans, while *Paper regions* include paragraphs, section headings, and list items.

They generate their corpus of 296 paper-presentation pairs from workshops of technical conferences, through simple searching. The papers were PDF format, and presentation were a mixture of PDF and PPT. Before working with them, they convert them into custom XML formats, which represent relevant parts of the original data as logical regions (orthographic boundaries). They prefer such physical regions over semantic regions for its simplicity to implement and verify.

The alignment problem now reduces to an IR query, where the query is a slide region (which is the section/subsection the slide is related to), and the documents are the target regions from the paper. They compare two TF-IDF based scoring methods, with and without query expansion, resulting in 4 alignment methods.

The procedure of scoring is as follows. For each token in each region TF-IDF is computed, where TF is the frequency of the token in the region and DF is the number of regions containing the token's stem. The slide region is tokenized and POS tagged to remove non-content words. Each token in the query is stemmed and then may or may not be query-expanded depending on the method. A score is then calculated for each target region with the query. Two scoring methods were used - one uses the average TF-IDF score of the search terms relative to the target region, and the other uses the quantity of the matched terms.

The evaluation of their methods gives critical insights. First, a vast majority of the slide regions are not alignable (zero score with all target regions) - meaning that a lot of information in slides is not present in the paper - contrary to their hypothesis. Then they define an *alignable* accuracy against the *raw* accuracy, considering only alignable slide regions. They find that the best algorithm on an average gives an average 75% alignable accuracy, but only 50% raw accuracy. Query expansion seems to have little or negative impact on the aligners and that the second scoring method is better than the first.

After more results, they conclude that the data indicates that the task of presentation generation is highly dependent on the end purpose the presentation will serve, as well as the target audience and other factors. Also query expansion generally degraded performance, possibly because authors tend to use wording in slides similar to their paper, and that using synonyms for query expansion is not aggressive

enough, and may require hypernyms, immediate hyponyms, and other semantically related terms. Also a possible loss of accuracy could be polysemy of words, causing query expansion to be incorrect and insensitive to context.

4. Approach

Our approach consists of 3 stages, parsing, summarization and slides generation. Parsing converts the HTML into a clean, well-defined XML format that is used for further processing. Summarization prunes the XML using a series of heuristics, to select the sections/lines that would be most relevant for the presentation. Finally, generator stage converts the pruned XML into L^AT_EX Beamer presentation. Each of the stages have an intermediate output, that can optionally be tweaked manually by user to achieve better results later, such as tweaking the set of keywords extracted after parsing stage can help summarization work better. We discuss each stage in detail next.

4.1. Parsing

Some cool stuff about parsing

4.2. Summarization

The second step in the slide generation process is that of summarization of the text as present in the complete research paper. The parsing step returns the structure of the paper in our predefined XML format, with all the lines, sections, subsections in the hierarchical order, and the objective of this stage is to select only those lines that are most informative, and would be most relevant for a presentation. We applied different heuristic techniques to summarize different parts of the paper, as discussed in following sections.

4.2.1 Summarizing Introduction

Introduction is an important section that introduces the problem, the motivation to solving it and gives an overall direction towards the solution, along with discussion an overview of the contributions. Abstract of the paper also summarizes the paper, albeit in a more succinct format. Hence, we use the approach as proposed in [2]. For each sentence in introduction, we compute it's TF-IDF based cosine similarity with the complete abstract text, and select the the top n sentences with highest scores.

4.2.2 Summarizing Model

We consider the central part of the paper, with the approach, results, analysis etc. to be the model part of the paper, that contains the bulk of information content. The information itself may be divided into multiple subsections and bulleted points, and may contain images, tables, mathematical proofs/expressions etc.

Based on our experiments, we observed that the sections tree structure as extracted by the parser is an important part of the paper, and we include each of the sections and subsections in the final set of slides in the same hierarchy as detected in original paper. However, we summarize the text content in each of those sections recursively. We use the following set of heuristics for summarizing lines in each:

1. **Bulleted Points** Some sections in the papers itself contain a set of points in bulleted fashion, or with numbering. Such points might also be spread across with text in between, but have continuous numbering. In such a case, we only take those points to be representative text for that section.
2. **Other Text** In case the section does not contain bullets, we manually select a subset of the lines from that section. We determine a score for every line by computing its TF-IDF score with the set of keywords relevant to the paper. Lines with frequent occurrence of keywords is usually more relevant, and we tend to select such sentences.




Figure 1. Block Diagram of the system

3. **Boosting image references** Since images are an important part of presentations, we boost the score of lines with image references. That increases the possibility of that line getting included, and hence of the associated image getting included in the slides as well.

Keyword Set Expansion The model section summarization depends heavily on the set of keywords, which we mostly extract from the keywords section of the paper itself. However, this set is usually very small, and sometimes ineffective for selection of best subset of sentences, hence we propose an algorithm to expand that set using the paper itself. We use co-occurrence probabilities to compute the other potential keywords, according to Algorithm 1.

Data: allLinesSet, keywordsSet

Result: expandedKeywordsSet

pairs = []

for line in lines **do**

if line contains any keyword **then**

for token in line.tokenized **do**

 pairs \leftarrow (token, keyword)

end

end

for pair in pairs **do**

if frequency > threshold **then**

 expandedKeywordsSet \leftarrow pair.token

end

end

end

Algorithm 1: Keyword Set Expansion

as [”proposed”, ”shown”, ”contributed”, concluded”] and so on, that are present in important sentences in conclusion. We compute TF-IDF based cosine similarity of each conclusion sentence with the above set, and select the top n .

4.2.4 Selecting images, references

We select graphics for display in the slides only if the sentence referring to that graphic gets selected. However, since images are an integral part of presentations, we boost the score of

4.3. Slides Generation

4.2.3 Summarizing Conclusion

Summarizing conclusion requires us to select sentences that mention the major contributions of the paper, along with its consequences and implications. Here we use the approach as proposed in [2]. We define a set of words such

5. Datasets

5.1. Chi96 Dataset Details

how procured? details etc

6. Experimental Results

what metric was used for evaluation, experimental settings, accuracy, plots (if any), tables.

7. Analysis of Results

8. Discussion

9. Conclusion

10. Future Work

- why you think intuitively one approach worked better than the other (if you tried two approaches), why do you think your approach works in some

References

- [1] B. Beamer and R. Girju. Investigating automatic alignment methods for slide generation from academic papers. In *Proceedings of the Thirteenth Conference on Computational Natural Language Learning*, CoNLL '09, pages 111–119,

Stroudsburg, PA, USA, 2009. Association for Computational Linguistics.

- [2] M. Sravanthi, C. R. Chowdary, and P. S. Kumar. Slidesgen: Automatic generation of presentation slides for a technical paper using summarization. In H. C. Lane and H. W. Guesgen, editors, *FLAIRS Conference*. AAAI Press, 2009.